



Root-cause checks: fitting is possible; reliable development gains remain unproven

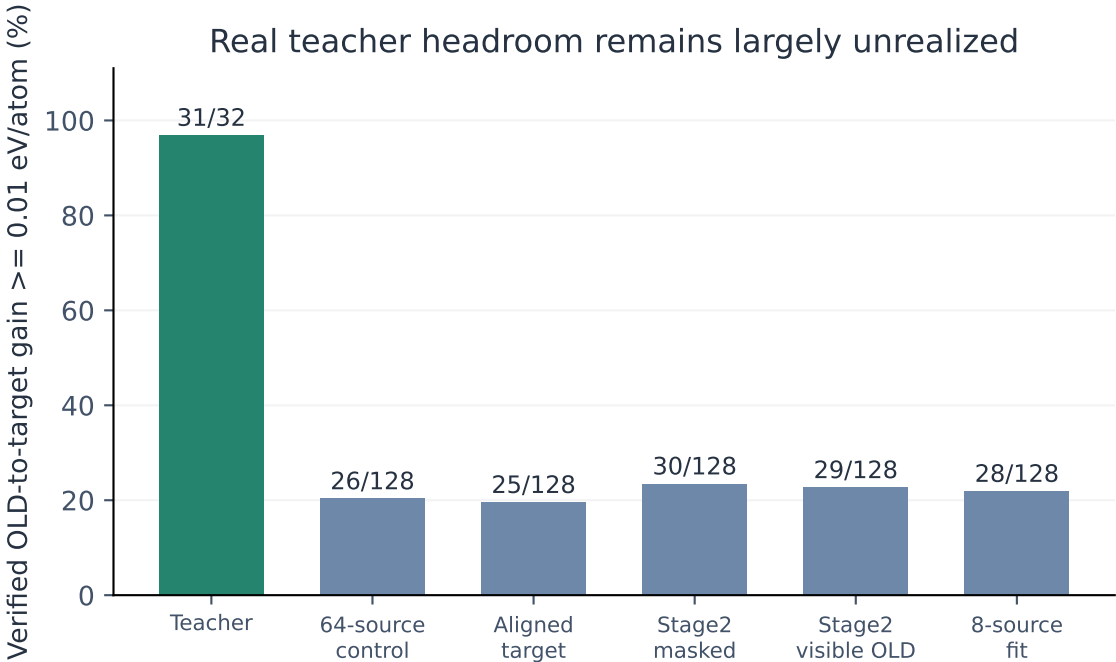
G: the tiny training set can be fitted



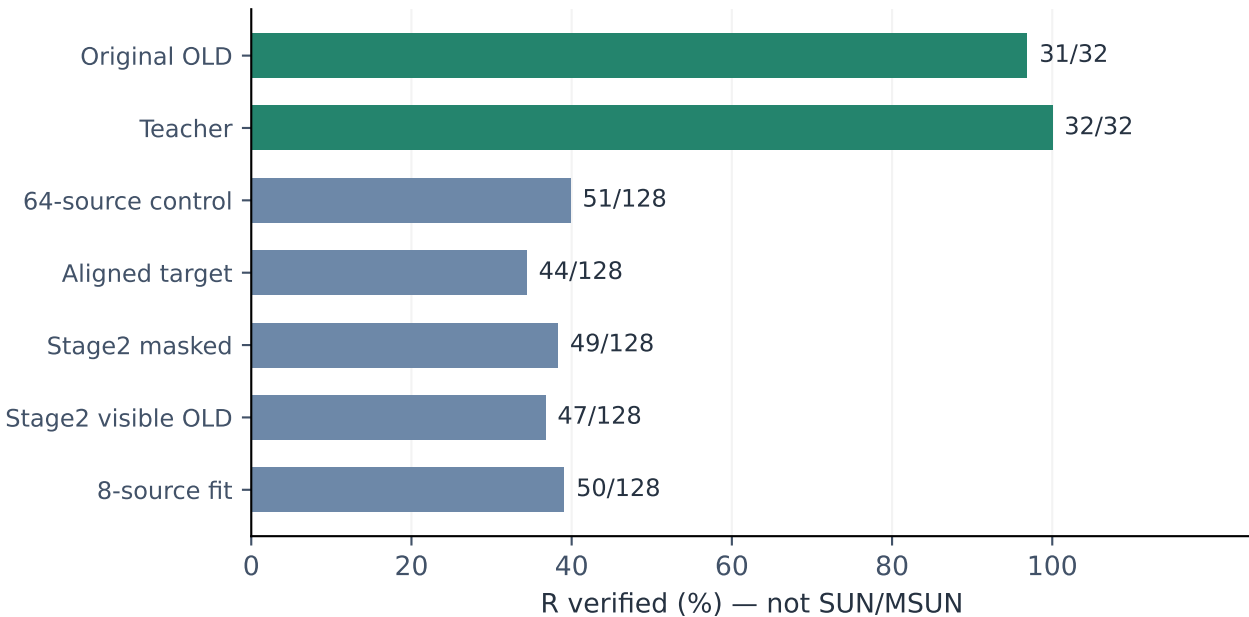
S: the tiny training set can be fitted



Real teacher headroom remains largely unrealized



S proposals do not preserve OLD reliability



Exploratory development diagnostics, not SUN/MSUN. Four student seeds reuse the same 32 S sources.
Counts across checkpoints are descriptive, not isolated causal effects. Paired A/B intervals are recorded in the companion data and report.